

```
In [2]: import numpy as np
from numpy import random
import seaborn as sns
import matplotlib.pyplot as plt
import tensorflow as tf
import pandas as pd
from tensorflow.keras.datasets import mnist
from tensorflow.keras.layers import Input, LSTM, Dense, GlobalMaxPool1D, GlobalMaxPool2D
from tensorflow.keras.models import Sequential
```

We start by importing the MNIST dataset and reshaping them.

```
In [3]: (X, y), (X_test, y_test) = mnist.load_data()
print('X_train: ' + str(X.shape))
print('Y_train: ' + str(y.shape))
print('X_test: ' + str(X_test.shape))
print('Y_test: ' + str(y_test.shape))
#pyplot.imshow(X[i], cmap=pyplot.get_cmap('gray')) gives image
```

```
X_train: (60000, 28, 28)
Y_train: (60000,)
X_test: (10000, 28, 28)
Y_test: (10000,)
```

```
In [4]: X_train = pd.DataFrame(X.reshape(-1, 784))
X_test = pd.DataFrame(X_test.reshape(-1, 784))
y_train = pd.DataFrame(y)[0]
y_test = pd.DataFrame(y_test)[0]
print('X_train: ' + str(X_train.shape))
print('Y_train: ' + str(y_train.shape))
print('X_test: ' + str(X_test.shape))
print('Y_test: ' + str(y_test.shape))
X_train, X_test = X_train / 255.0, X_test / 255.0
```

```
X_train: (60000, 784)
Y_train: (60000,)
X_test: (10000, 784)
Y_test: (10000,)
```

We write the functions for creation of **ANN** and other useful ones

```

In [4]: def create_ann(op=10):
    model = tf.keras.models.Sequential([
        tf.keras.layers.Input(shape=(features,)),
        tf.keras.layers.Dense(128, activation='relu'),
        tf.keras.layers.Dropout(0.2),
        tf.keras.layers.Dense(128, activation='relu'),
        tf.keras.layers.Dropout(0.2),
        tf.keras.layers.Dense(128, activation='relu'),
        tf.keras.layers.Dropout(0.2),
        tf.keras.layers.Dense(op, activation='softmax')
    ])
    model.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])
    return model

def confidence(model, ip, label=None):
    v= ip.values
    v=v.reshape(-1,1)
    if not label:
        a=list(model.predict(v.T)[0])
        return a.index(max(a)),max(a)
    return model.predict(v.T)[0][label]

def confidence_distribution(model, ip,ind = 0,op=10):
    try:
        model = model.model
    except:
        pass
    ip=ip.iloc[ind,:]
    v= ip.values
    v=v.reshape(-1,1)
    l=model.predict(v.T)[0]
    plt.bar([i for i in range(op)],1)
    plt.show()

def show_mistake(model,df,i=None,y_test = y_test):
    try:
        model = model.model
    except:
        pass
    p_test = model.predict(df).argmax(axis=1)
    wrong_idx = np.where(p_test != y_test)[0]
    if i is None:
        i = np.random.choice(wrong_idx)
    display(df, y_test, i)
    print("Predicted: ",p_test[i])
    return i

def display(df, ydf, index):
    dd = df.iloc[index,:].to_numpy() *255
    try:
        if not ydf:

```

```

        plt.imshow(dd.reshape(28,28), cmap='gray')
    except:
        plt.title("True label: %s" % (ydf[index]))
        plt.imshow(dd.reshape(28,28), cmap='gray')

def threshold(df,t_val):
    out = df.copy()
    for i in range(df.shape[0]):
        out.iloc[i,:] = df.iloc[i,:].apply(lambda x: 0 if x<t_val else 1)
    return out

def model_stats(model):
    a_float = model.history['val_accuracy'][-1]*100
    plt.plot(model.history['accuracy'], label='train_acc')
    plt.plot(model.history['val_accuracy'], label='val_acc')
    plt.title("{:.2f}".format(a_float)+"% accuracy")
    plt.legend()

def load(m):
    try:
        return tf.keras.models.load_model(m)
    except OSError:
        print(str(m)+" not found")

def noise_analysis(model, conf_th = 0.75, noise_n = 100):
    noise = pd.DataFrame(np.random.rand(noise_n, 784))
    nums=[]
    cnum=[]
    try:
        model = model.model
    except:
        pass
    for i in range(noise.shape[0]):
        num,conf = confidence(model, noise.iloc[i,:])
        nums.append(num)
        if conf>conf_th:
            cnum.append(num)
    plt.hist(nums)
    plt.xticks(range(11))
    plt.title("Distribution of "+str(len(nums))+" noise data sets and their predi")
    plt.show()
    plt.hist(cnum)
    plt.xticks(range(11))
    plt.title("Distribution of "+str(len(cnum))+"/"+str(len(nums))+" noise data s")
    plt.show()


def shuff_together(df1,df2):
    joined = pd.concat([df1,df2], axis=1)
    joined = joined.iloc[np.random.permutation(len(joined))].reset_index(drop=True)
    return joined.iloc[:, :-1],joined.iloc[:, -1]

def model_predictions(model,df,y_test=y_test):
    try:
        model = model.model

```

```
except:
    pass
preds = model.predict(df)
l=[list(preds[i]).index(max(preds[i])) for i in range(preds.shape[0])]
plt.hist(y_test, label=['Actual values'])
plt.hist(l, label=['Predicted values'])
plt.legend(loc='upper right')
plt.title("Distribution of model predicted values vs actual labels")
plt.show()
return l

display(X_test, y_test, 8097)
```




Standard MNIST **training** using ANN


```
In [5]: features = X_train.shape[1]

model = create_ann()
ann = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=10, verbose=1)
ann.model.save("ann")
model_stats(ann)
#Load("ann")
```

INFO:tensorflow:Assets written to: ann/assets




```
In [6]: _ = model_predictions(ann,X_test)
```




We look at one of the wrongly labelled test cases as well as its probability distribution prediction by the model.

```
In [ ]: #Shows random mistake each time this cell is executed  
i=show_mistake(ann,X_test)
```

Predicted: 6



```
In [ ]: confidence_distribution(ann, X_test,i)
```




Intelligence of a model is how well it learns to do its job correctly given a dataset. It is inversely proportional to the size of the training set.

Here we will limit the training set to *maxrows*


In []: *#INTELLIGENCE*

```
model = create_ann()  
maxrows = 1000  
r = model.fit(X_train.iloc[:maxrows,:], y_train.iloc[:maxrows:], validation_data=  
model_stats(r)
```




Trying pre processing of data using a threshold value. Below, everything above 90% of the pixel value is changed to 100% while other pixels are at 0%


```
In [ ]: X_train_50 = threshold(X_train,0.9)
X_test_50 = threshold(X_test,0.9)
display(X_test_50, y_test, 50)
```




This does not seem to increase the intelligence when we test on normal test as well as pre processed test data

```
In [ ]: model = create_ann()
r = model.fit(X_train_50.iloc[:maxrows,:], y_train.iloc[:maxrows:], validation_data=(X_test_50, y_test))
model_stats(r)
```




```
In [ ]: model = create_ann()  
r = model.fit(X_train_50.iloc[:maxrows,:], y_train.iloc[:maxrows:], validation_data=(X_test, y_test))  
model_stats(r)
```



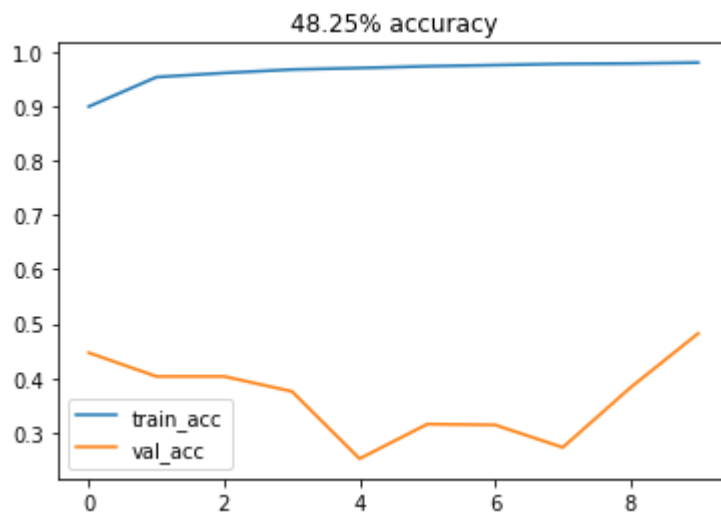
We will create **noisy** test data

```
In [ ]: X_test_noisy = X_test.apply(lambda x: x+random.rand())  
display(X_test_noisy, y_test, 8097)
```




Trying to see if this model **ANN** works on this as a test of versatility and robustness

```
In [ ]: model = create_ann()
ann = model.fit(X_train, y_train, validation_data=(X_test_noisy, y_test), epochs=10)
#ann.model.save("ann")
model_stats(ann)
```




We will try to **train with noisy data** and see

```
In [ ]: X_train_noisy = X_train.apply(lambda x: x+random.rand())
model = create_ann()
ann_noisy = model.fit(X_train_noisy, y_train, validation_data=(X_test_noisy, y_test), epochs=10)
model_stats(ann_noisy)
```




```
In [ ]: _ = model_predictions(ann_noisy,X_test_noisy)
```




```
In [ ]: i=show_mistake(ann_noisy,X_test_noisy)
```

Predicted: 2




```
In [ ]: confidence_distribution(ann_noisy, X_test_noisy,i)
```




The model trained with noisy data (ann_noisy) seems to give much better accuracy for noisy test data and as good accuracy for the normal test data shown below making it more versatile

```
In [ ]: model = create_ann()
ann_noisy = model.fit(X_train_noisy, y_train, validation_data=(X_test, y_test), epochs=10)
model_stats(ann_noisy)
ann_noisy.model.save("ann_noisy")
```

INFO:tensorflow:Assets written to: ann_noisy\assets




```
In [ ]: _ = model_predictions(ann_noisy,X_test)
```




```
In [ ]: i=show_mistake(ann_noisy,X_test)
```

Predicted: 7




```
In [ ]: confidence_distribution(ann_noisy, X_test,i)
```



Creating a **pure noise** dataset made of random pixels

```
In [ ]: noise = pd.DataFrame(np.random.rand(100, 784))  
display(noise, None, 55)
```



ann_noisy is more like to label a pure noise data with high level of certainty

```
In [ ]: ind = 32
print("ann_noisy model's label guess and confidence on a pure noise image is ",confidence[ind])
print("ann model's label guess and confidence on a pure noise image is ",confidence[ind])
```


ann_noisy model's label guess and confidence on a pure noise image is (8, 0.9446118)

ann model's label guess and confidence on a pure noise image is (3, 0.9991881)


ann model is used to predict a set of noisy images. We find that a small number of them is labeled with very high confidence. It does not have high confidence in its predictions for noise data

```
In [ ]: noise_analysis(load('ann'),0.9)
```

Distribution of 100 noise data sets and their predicted labels




Distribution of 87/100 noise data sets that show a confidence above 90.0%




ann_noisy model is used to predict a set of noisy images. We find that fewer of them are labeled with very high confidence


```
In [ ]: noise_analysis(ann_noisy,0.9)
```

Distribution of 100 noise data sets and their predicted labels



Distribution of 37/100 noise data sets that show a confidence above 90.0%



We append noise dataset to our original data and assign them to label:10


```

In [ ]: noise = pd.DataFrame(np.random.rand(5000, 784))
X_train_plusnoisy = pd.concat([X_train, noise], axis=0)
y_min1 = pd.DataFrame(np.full((5000, ), 10))
y_train_plusnoisy = pd.concat([y_train, y_min1], axis=0)[0]
X_train_plusnoisy, y_train_plusnoisy = shuff_together(X_train_plusnoisy, y_train_plusnoisy)
display(X_train_plusnoisy, y_train_plusnoisy, y_train_plusnoisy[y_train_plusnoisy =

noise = pd.DataFrame(np.random.rand(1000, 784))
X_test_plusnoisy = pd.concat([X_test, noise], axis=0)
y_min1 = pd.DataFrame(np.full((1000, ), 10))
y_test_plusnoisy = pd.concat([y_test, y_min1], axis=0)[0]
X_test_plusnoisy, y_test_plusnoisy = shuff_together(X_test_plusnoisy, y_test_plusnoisy)
X_train_plusnoisy.shape, y_train_plusnoisy.shape

```

Out[28]: ((65000, 784), (65000,))




We now train an ANN with pure noise set as a separate input

```
In [ ]: features = X_train_plusnoisy.shape[1]

model = create_ann(11)
ann_plusnoisy = model.fit(X_train_plusnoisy, y_train_plusnoisy, validation_data=
ann_plusnoisy.model.save("ann_plusnoisy")
model_stats(ann_plusnoisy)
#Load("ann")
```


INFO:tensorflow:Assets written to: ann_plusnoisy\assets




This classifies noise correctly

```
In [ ]: noise_analysis(ann_plusnoisy,0.9)
```

Distribution of 100 noise data sets and their predicted labels




Distribution of 100/100 noise data sets that show a confidence above 90.0%



We will create a training set that is noisy data plus noise

```
In [ ]: noise = pd.DataFrame(np.random.rand(5000, 784))
X_trainnoisy_plusnoise = pd.concat([X_train_noisy, noise], axis=0)
y_min1 = pd.DataFrame(np.full((5000, ), 10))
y_trainnoisy_plusnoise = pd.concat([y_train, y_min1], axis=0)[0]
X_trainnoisy_plusnoise, y_trainnoisy_plusnoise = shuff_together(X_train_plusnoisy,
display(X_trainnoisy_plusnoise, y_trainnoisy_plusnoise, y_trainnoisy_plusnoise[y_tr
```




We can create an ANN that is trained on noisy data and noise and it does a good job of labelling original data and noise correctly

```
In [ ]: features = X_trainnoisy_plusnoise.shape[1]

model = create_ann(11)
ann_noisy_plusnoise = model.fit(X_trainnoisy_plusnoise, y_trainnoisy_plusnoise, v
ann_noisy_plusnoise.model.save("ann_noisy_plusnoise")
model_stats(ann_noisy_plusnoise)
#Load("ann")
```

INFO:tensorflow:Assets written to: ann_noisy_plusnoise/assets




But it does a bad job of labelling noisy test data with most predictions being 5

```
In [ ]: ann_noisy_plusnoise.model.evaluate(X_test_noisy,y_test)
```

```
313/313 [=====] - 1s 1ms/step - loss: 7.8381 - accurac  
y: 0.1930
```

```
Out[33]: [7.838090419769287, 0.19300000369548798]
```


```
In [ ]: _ = model_predictions(ann_noisy_plusnoise,X_test_noisy)
```




This also seem to show very high confidence in its incorrectly predicted labels.

```
In [ ]: i=show_mistake(ann_noisy_plusnoise,X_test_noisy)
```

Predicted: 5




```
In [ ]: confidence_distribution(ann_noisy_plusnoise, X_test_noisy,i,op=11)
```




What will happen if the model is faced with inverted data?

```
In [ ]: inverted = X_test.apply(lambda x: 1-x)  
display(inverted,y_test,ind)
```




Every thing is labelled as noise with high certainty

```
In [ ]: model = create_ann(11)
r = model.fit(X_train_plusnoisy, y_train_plusnoisy, validation_data=(inverted, y_
#ann.model.save("ann")
model_stats(r)
```




```
In [ ]: _ = model_predictions(r,inverted)
```





```
In [ ]: i=show_mistake(r,inverted)
```

Predicted: 10




```
In [ ]: confidence_distribution(r, inverted,i,11)
```




The original standard ANN without noise label predicted most inverted images as 5 with very high confidence

```
In [ ]: model = create_ann()  
r = model.fit(X_train, y_train, validation_data=(inverted, y_test), epochs=10, ver  
#ann.model.save("ann")  
model_stats(r)
```




```
In [ ]: _ = model_predictions(r,inverted)
```




```
In [ ]: i=show_mistake(r,inverted)
```

Predicted: 5



```
In [ ]: confidence_distribution(r, inverted,i)
```



We will now create a model based on RNN

```

In [ ]: (X, y), (X_t, y_t) = mnist.load_data()
X, X_t = X / 255.0, X_t / 255.0

def create_rnn():
    i = Input(shape=X[0].shape)
    x = LSTM(128, return_sequences=True, dropout=0.2)(i)
    x = GlobalMaxPool1D()(x)
    x = Dense(10, activation='softmax')(x)

    model = tf.keras.models.Model(i, x)

    model.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])
    return model

def confidence(model, ip):
    try:
        model = model.model
    except:
        pass
    a=list(model.predict(np.expand_dims(ip, axis=0))[0])
    return a.index(max(a)),max(a)

def confidence_distribution(model, ip, ind = 0):
    try:
        model = model.model
    except:
        pass
    l=model.predict(ip.values.reshape(-1,1,28,28)[ind])[0]
    plt.bar([i for i in range(10)],l)
    plt.show()

def noise_analysis(model, conf_th = 0.75, noise_n = 100):
    noise = np.random.rand(noise_n, 28,28)
    nums=[]
    cnum=[]
    try:
        model = model.model
    except:
        pass
    for i in range(noise.shape[0]):
        num,conf = confidence(model, noise[i])
        nums.append(num)
        if conf>conf_th:
            cnum.append(num)
    plt.hist(nums)
    plt.xticks(range(11))
    plt.title("Distribution of "+str(len(nums))+" noise data sets and their predi")
    plt.show()
    plt.hist(cnum)
    plt.xticks(range(11))
    plt.title("Distribution of "+str(len(cnum))+"/"+str(len(nums))+" noise data s")
    plt.show()

```

```
def show_mistake(model,df,i=None,y_test = y_test):
    try:
        model = model.model
    except:
        pass
    try:
        p_test = model.predict(df.values.reshape(-1,28,28)).argmax(axis=1)
    except:
        p_test = model.predict(df.argmax(axis=1))
    wrong_idx = np.where(p_test != y_test)[0]
    if i is None:
        i = np.random.choice(wrong_idx)
    display(df, y_test, i)
    print("Predicted: ",p_test[i])
    return i

def model_predictions(model,df,y_test=y_test):
    try:
        model = model.model
    except:
        pass
    try:
        preds = model.predict(df.values.reshape(-1,28,28)).argmax(axis=1)
    except:
        preds = model.predict(df.argmax(axis=1))
    plt.hist(y_test, label=['Actual values'])
    plt.hist(preds, label=['Predicted values'])
    plt.legend(loc='upper right')
    plt.title("Distribution of model predicted values vs actual labels")
    plt.show()
    return preds
```

We have high accuracy similar to ANN

```
In [ ]: model = create_rnn()
rnn = model.fit(X, y, validation_data=(X_t, y_t), epochs=10, verbose=1)
model_stats(rnn)
```

Epoch 1/10

1875/1875 [=====] - 26s 13ms/step - loss: 0.7777 - accuracy: 0.7581 - val_loss: 0.1553 - val_accuracy: 0.9539

Epoch 2/10

1875/1875 [=====] - 24s 13ms/step - loss: 0.1461 - accuracy: 0.9564 - val_loss: 0.0827 - val_accuracy: 0.9754

Epoch 3/10

1875/1875 [=====] - 24s 13ms/step - loss: 0.0957 - accuracy: 0.9714 - val_loss: 0.0633 - val_accuracy: 0.9823

Epoch 4/10

1875/1875 [=====] - 24s 13ms/step - loss: 0.0762 - accuracy: 0.9774 - val_loss: 0.0492 - val_accuracy: 0.9844

Epoch 5/10

1875/1875 [=====] - 24s 13ms/step - loss: 0.0612 - accuracy: 0.9812 - val_loss: 0.0549 - val_accuracy: 0.9830


Epoch 6/10

1875/1875 [=====] - 24s 13ms/step - loss: 0.0489 - accuracy: 0.9842 - val_loss: 0.0369 - val_accuracy: 0.9882

Epoch 7/10


1875/1875 [=====] - 24s 13ms/step - loss: 0.0327 - accuracy: 0.9912 - val_loss: 0.0227 - val_accuracy: 0.9927

```
In [ ]: _ = model_predictions(rnn,X_test)
```




```
In [ ]: i=show_mistake(rnn,X_test)
```

Predicted: 2



```
In [ ]: confidence_distribution(rnn,X_test,i)
```




But its performance on noisy data is worse than ANN

```
In [ ]: rnn.model.evaluate(X_test_noisy.values.reshape(-1,28,28), y_test)
```

```
313/313 [=====] - 2s 5ms/step - loss: 5.4675 - accurac  
y: 0.1162
```


```
Out[51]: [5.467532157897949, 0.11620000004768372]
```

```
In [ ]: _ = model_predictions(rnn,X_test_noisy)
```





```
In [ ]: noise_analysis(rnn,0.9)
```

Distribution of 100 noise data sets and their predicted labels



Distribution of 24/100 noise data sets that show a confidence above 90.0%




```
In [ ]: noise_28by28 = noise.values.reshape(-1,28,28)
noise_28by28.shape
```

```
Out[54]: (5000, 28, 28)
```


The rnn predicts each noise as almost entirely 0

```
In [ ]: confidence_distribution(rnn,noise,np.random.randint(150))
```



Training and validation using noisy data

```
In [ ]: model = create_rnn()
rnn_noisy = model.fit(X_train_noisy.values.reshape(-1,28,28), y_train, validation_data=(X_test_noisy.values.reshape(-1,28,28), y_test))
model_stats(rnn_noisy)
```




This model does not do well in predicting the standard test data, labelling most of them as certain 2 numbers

```
In [ ]: rnn_noisy.model.evaluate(X_t, y_t)
```

```
313/313 [=====] - 2s 5ms/step - loss: 1.5631 - accuracy: 0.5597
```


```
Out[57]: [1.563126564025879, 0.5597000122070312]
```

```
In [ ]: _ = model_predictions(rnn_noisy,X_test)
```




```
In [ ]: i=show_mistake(rnn_noisy,X_test)
```

Predicted: 3




```
In [ ]: confidence_distribution(rnn_noisy,X_test,i)
```




But it does seem to recognize pure noise well, not being confident about them

```
In [ ]: noise_analysis(rnn_noisy,0.9)
```

Distribution of 100 noise data sets and their predicted labels




Distribution of 3/100 noise data sets that show a confidence above 90.0%



confidence_distribution of random noise data

```
In [ ]: confidence_distribution(rnn_noisy,noise,np.random.randint(150))
```



As expected, it does poorly with inverted data


```
In [ ]: rnn_noisy.model.evaluate(inverted.values.reshape(-1,28,28), y_t)
```

313/313 [=====] - 2s 5ms/step - loss: 5.9361 - accuracy: 0.0936


Out[63]: [5.9360527992248535, 0.09359999746084213]

The Intelligence seems to be comparable to the others

```
In [ ]: model = create_rnn()
rnn = model.fit(X[:maxrows,:,:), y[:maxrows:], validation_data=(X_test.values.reshape(-1,28,28), y_test), verbose=1)
model_stats(rnn)
```




```
In [ ]: _ = model_predictions(rnn,X_test)
```




The Intelligence of threshold trained model set seems to be similar

```
In [ ]: model = create_rnn()
rnn = model.fit(X_train_50[:maxrows].values.reshape(-1,28,28), y[:maxrows,], val_i
model_stats(rnn)
```



```
In [ ]: _ = model_predictions(rnn,X_test_50)
```



We will now create a model based on CNN

```

In [21]: (X, y), (X_t, y_t) = mnist.load_data()

def create_cnn(op=10):
    model = tf.keras.models.Sequential([
        tf.keras.layers.Input(shape=(28, 28, 1)),
        tf.keras.layers.Conv2D(256, (3, 3), activation='relu'),
        tf.keras.layers.Dropout(0.2),
        tf.keras.layers.MaxPooling2D(pool_size=(2, 2)),
        #tf.keras.layers.Dropout(0.2),
        tf.keras.layers.Flatten(),
        tf.keras.layers.Dense(64),
        tf.keras.layers.Dropout(0.2),
        tf.keras.layers.Dense(op, activation='softmax')])
    model.compile(loss='sparse_categorical_crossentropy',
                  optimizer='adam',
                  metrics=['accuracy'])
    return model

def model_stats(model):
    a_float = model.history.history['val_accuracy'][-1]*100
    plt.plot(model.history.history['accuracy'], label='train_acc')
    plt.plot(model.history.history['val_accuracy'], label='val_acc')
    plt.title("{:.2f}".format(a_float)+"% accuracy")
    plt.legend()

def model_predictions(model, df, y_test=y_test):
    try:
        model = model.model
    except:
        pass
    #preds = model.predict(df)
    try:
        l=model.predict(df).argmax(axis=1)
    except:
        l=model.predict(df.values.reshape(-1,28, 28,1)).argmax(axis=1)
    #l=[list(preds[i]).index(max(preds[i])) for i in range(preds.shape[0])]
    plt.hist(y_test, label=['Actual values'])
    plt.hist(l, label=['Predicted values'])
    plt.legend(loc='upper right')
    plt.title("Distribution of model predicted values vs actual labels")
    plt.show()
    return l

def show_mistake(model, df, i=None, y_test = y_test):
    try:
        model = model.model
    except:
        pass
    try:
        p_test=model.predict(df).argmax(axis=1)
    except:
        p_test=model.predict(df.values.reshape(-1,28, 28,1)).argmax(axis=1)
    #p_test = model.predict(df).argmax(axis=1)
    wrong_idx = np.where(p_test != y_test)[0]
    if i is None:
        i = np.random.choice(wrong_idx)

```



```

display(df, y_test, i)
print("Predicted: ",p_test[i])
return i

def confidence_distribution(model, df,ind = 0,op=10):
    try:
        model = model.model
    except:
        pass
    try:
        l=model.predict(df[ind])[0]
    except:
        ip = df.values.reshape(-1,28, 28,1)
        l=model.predict(ip[ind].reshape(-1,28, 28,1))[0]
    plt.bar([i for i in range(op)],1)
    plt.show()

def confidence(model, ip, label=None):
    try:
        model = model.model
    except:
        pass
    #a=model.predict(ip)[0]
    try:
        a=model.predict(ip)[0]
    except:
        ip = ip.values.reshape(-1,28, 28,1)
        a=model.predict(ip.reshape(-1,28, 28,1))[0]
    #v= ip.values
    #v=v.reshape(-1,1)
    a=list(a)
    if not label:
        return a.index(max(a)),max(a)
    return a[label]


def noise_analysis(model, conf_th = 0.75, noise_n = 100):
    #noise = pd.DataFrame(np.random.rand(noise_n, 784))
    noise = np.random.rand(100, 1,28,28,1)
    nums=[]
    cnum=[]
    try:
        model = model.model
    except:
        pass
    for i in range(noise.shape[0]):
        num,conf = confidence(model, noise[i])
        nums.append(num)
        if conf>conf_th:
            cnum.append(num)
    plt.hist(nums)
    plt.xticks(range(11))
    plt.title("Distribution of "+str(len(nums))+" noise data sets and their predi")
    plt.show()
    plt.hist(cnum)
    plt.xticks(range(11))

```

```
plt.title("Distribution of "+str(len(cnum))+"/"+str(len(nums))+" noise data s")
plt.show()
```


Each data point accepted by the CNN must be 3 dimensional

```
In [ ]: cnn = create_cnn()
#X = X.reshape(-1,28, 28,1)
#X_t = X_t.reshape(-1,28, 28,1)
cnn.fit(np.array(X_train).reshape(-1,28, 28,1), y, epochs=2, validation_data=(np.
model_stats(cnn)
```




```
In [ ]: i=show_mistake(cnn,X_test)
```

Predicted: 8




```
In [ ]: confidence_distribution(cnn, X_test,i)
```



Intelligence

```
In [22]: maxrows = 1000

cnn = create_cnn()
cnn.fit(np.array(X_train).reshape(-1,28, 28,1)[:maxrows,:,:), y[:maxrows], epochs
model_stats(cnn)
```




Trying to see if this CNN model works on noisy data, results appear worse than for ANN, with most predictions showing 1,7

```
In [ ]: cnn.evaluate(np.array(X_test_noisy).reshape(-1,28, 28,1),y_t)
```

313/313 [=====] - 4s 12ms/step - loss: 68.1384 - accur
acy: 0.2126


Out[72]: [68.13839721679688, 0.212599927520752]

```
In [ ]: _ = model_predictions(cnn,X_test_noisy)
```



This is how it views pure noise data


```
In [ ]: noise = np.random.rand(100, 1,28,28,1)
confidence_distribution(cnn, noise,10)
```




It does seem to recognize pure noise well, not being confident about them

```
In [ ]: noise_analysis(cnn,0.9)
```

Distribution of 100 noise data sets and their predicted labels



Distribution of 0/100 noise data sets that show a confidence above 90.0%




We see much better results when faced with inverted test data

```
In [ ]: cnn.evaluate(np.array(inverted).reshape(-1,28, 28,1),y_t)
```

```
313/313 [=====] - 4s 12ms/step - loss: 1.2094 - accuracy: 0.6287
```


```
Out[76]: [1.2093627452850342, 0.6287000179290771]
```

```
In [ ]: _ = model_predictions(cnn,inverted)
```



Training CNN using noisy train data

```
In [ ]: cnn = create_cnn()
cnn.fit(np.array(X_train_noisy).reshape(-1,28, 28,1), y, epochs=2, validation_data=(X_test, y))
model_stats(cnn)
```




When trained on noisy data, CNN gives above 90% accuracy for both noisy and original test data

```
In [ ]: cnn.evaluate(np.array(X_test).reshape(-1,28, 28,1),y_t)
```

313/313 [=====] - 4s 12ms/step - loss: 0.2029 - accuracy: 0.9351

Out[79]: [0.20290842652320862, 0.9351000189781189]

```
In [ ]: _ = model_predictions(cnn,X_test)
```



But gives worse performance on inverted test data

```
In [ ]: cnn.evaluate(np.array(inverted).reshape(-1,28, 28,1),y_t)
```

313/313 [=====] - 4s 12ms/step - loss: 5.8447 - accuracy: 0.3773

Out[81]: [5.844680309295654, 0.3772999942302704]

```
In [ ]: _ = model_predictions(cnn,np.array(inverted).reshape(-1,28, 28,1))
```

